

# untitled7

September 28, 2023

```
[1]: '''1.Download the Employee Attrition Dataset
https://www.kaggle.com/datasets/patelprashant/employee-attrition
2.Perfrom Data Preprocessing
3.Model Building using Logistic Regression and Decision Tree and Random Forest
4.Calculate Performance metrics'''
```

```
[1]: '1.Download the Employee Attrition
Dataset\nhttps://www.kaggle.com/datasets/patelprashant/employee-
attrition\n2.Perfrom Data Preprocessing\n3.Model Building using Logistic
Regression and Decision Tree and Random Forest\n4.Calculate Performance metrics'
```

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[3]: df=pd.read_csv('Employee-Attrition.csv')
df
```

```
[3]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	\
0	41	Yes	Travel_Rarely	1102		Sales
1	49	No	Travel_Frequently	279	Research & Development	
2	37	Yes	Travel_Rarely	1373	Research & Development	
3	33	No	Travel_Frequently	1392	Research & Development	
4	27	No	Travel_Rarely	591	Research & Development	
...	...	...	...	...	...	...
1465	36	No	Travel_Frequently	884	Research & Development	
1466	39	No	Travel_Rarely	613	Research & Development	
1467	27	No	Travel_Rarely	155	Research & Development	
1468	49	No	Travel_Frequently	1023		Sales
1469	34	No	Travel_Rarely	628	Research & Development	

	DistanceFromHome	Education	EducationField	EmployeeCount	\
0		1	2 Life Sciences		1
1		8	1 Life Sciences		1
2		2	2 Other		1
3		3	4 Life Sciences		1

4	2	1	Medical	1
...	...	...	...	...
1465	23	2	Medical	1
1466	6	1	Medical	1
1467	4	3	Life Sciences	1
1468	2	3	Medical	1
1469	8	3	Medical	1

	EmployeeNumber	...	RelationshipSatisfaction	StandardHours	\
0	1	...	1	80	
1	2	...	4	80	
2	4	...	2	80	
3	5	...	3	80	
4	7	...	4	80	
...	...	...	...	...	
1465	2061	...	3	80	
1466	2062	...	1	80	
1467	2064	...	2	80	
1468	2065	...	4	80	
1469	2068	...	1	80	

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
0	0	8	0	
1	1	10	3	
2	0	7	3	
3	0	8	3	
4	1	6	3	
...	...	...	...	
1465	1	17	3	
1466	1	9	5	
1467	1	6	0	
1468	0	17	3	
1469	0	6	3	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
0	1	6	4	
1	3	10	7	
2	3	0	0	
3	3	8	7	
4	3	2	2	
...	...	...	...	
1465	3	5	2	
1466	3	7	7	
1467	3	6	2	
1468	2	9	6	
1469	4	4	3	

	YearsSinceLastPromotion	YearsWithCurrManager
0	0	5
1	1	7
2	0	0
3	3	0
4	2	2
...	...	...
1465	0	3
1466	1	7
1467	0	3
1468	0	8
1469	1	2

[1470 rows x 35 columns]

```
[4]: df.isnull().sum()
```

```
[4]: Age                                0
Attrition                             0
BusinessTravel                         0
DailyRate                             0
Department                             0
DistanceFromHome                       0
Education                              0
EducationField                          0
EmployeeCount                           0
EmployeeNumber                          0
EnvironmentSatisfaction                 0
Gender                                  0
HourlyRate                             0
JobInvolvement                         0
JobLevel                               0
JobRole                                 0
JobSatisfaction                         0
MaritalStatus                           0
MonthlyIncome                           0
MonthlyRate                             0
NumCompaniesWorked                     0
Over18                                  0
OverTime                                0
PercentSalaryHike                       0
PerformanceRating                       0
RelationshipSatisfaction                 0
StandardHours                           0
StockOptionLevel                        0
TotalWorkingYears                       0
TrainingTimesLastYear                   0
```

```

WorkLifeBalance      0
YearsAtCompany        0
YearsInCurrentRole    0
YearsSinceLastPromotion 0
YearsWithCurrManager  0
dtype: int64

```

```
[5]: df.head()
```

```

[5]:   Age Attrition      BusinessTravel  DailyRate      Department \
0   41      Yes      Travel_Rarely      1102      Sales
1   49      No  Travel_Frequently      279  Research & Development
2   37      Yes      Travel_Rarely      1373  Research & Development
3   33      No  Travel_Frequently      1392  Research & Development
4   27      No      Travel_Rarely      591  Research & Development

      DistanceFromHome  Education EducationField  EmployeeCount  EmployeeNumber \
0                1      2  Life Sciences      1      1
1                8      1  Life Sciences      1      2
2                2      2      Other      1      4
3                3      4  Life Sciences      1      5
4                2      1      Medical      1      7

      ... RelationshipSatisfaction  StandardHours  StockOptionLevel \
0      ...                1      80      0
1      ...                4      80      1
2      ...                2      80      0
3      ...                3      80      0
4      ...                4      80      1

      TotalWorkingYears  TrainingTimesLastYear  WorkLifeBalance  YearsAtCompany \
0                8      0      1      6
1               10      3      3     10
2                7      3      3      0
3                8      3      3      8
4                6      3      3      2

      YearsInCurrentRole  YearsSinceLastPromotion  YearsWithCurrManager
0                4      0      5
1                7      1      7
2                0      0      0
3                7      3      0
4                2      2      2

```

```
[5 rows x 35 columns]
```

```
[6]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                            1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                             1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                        1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                       1470 non-null   int64
17  MaritalStatus                        1470 non-null   object
18  MonthlyIncome                        1470 non-null   int64
19  MonthlyRate                          1470 non-null   int64
20  NumCompaniesWorked                   1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                    1470 non-null   int64
24  PerformanceRating                    1470 non-null   int64
25  RelationshipSatisfaction              1470 non-null   int64
26  StandardHours                        1470 non-null   int64
27  StockOptionLevel                     1470 non-null   int64
28  TotalWorkingYears                    1470 non-null   int64
29  TrainingTimesLastYear                1470 non-null   int64
30  WorkLifeBalance                      1470 non-null   int64
31  YearsAtCompany                       1470 non-null   int64
32  YearsInCurrentRole                   1470 non-null   int64
33  YearsSinceLastPromotion               1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB

```

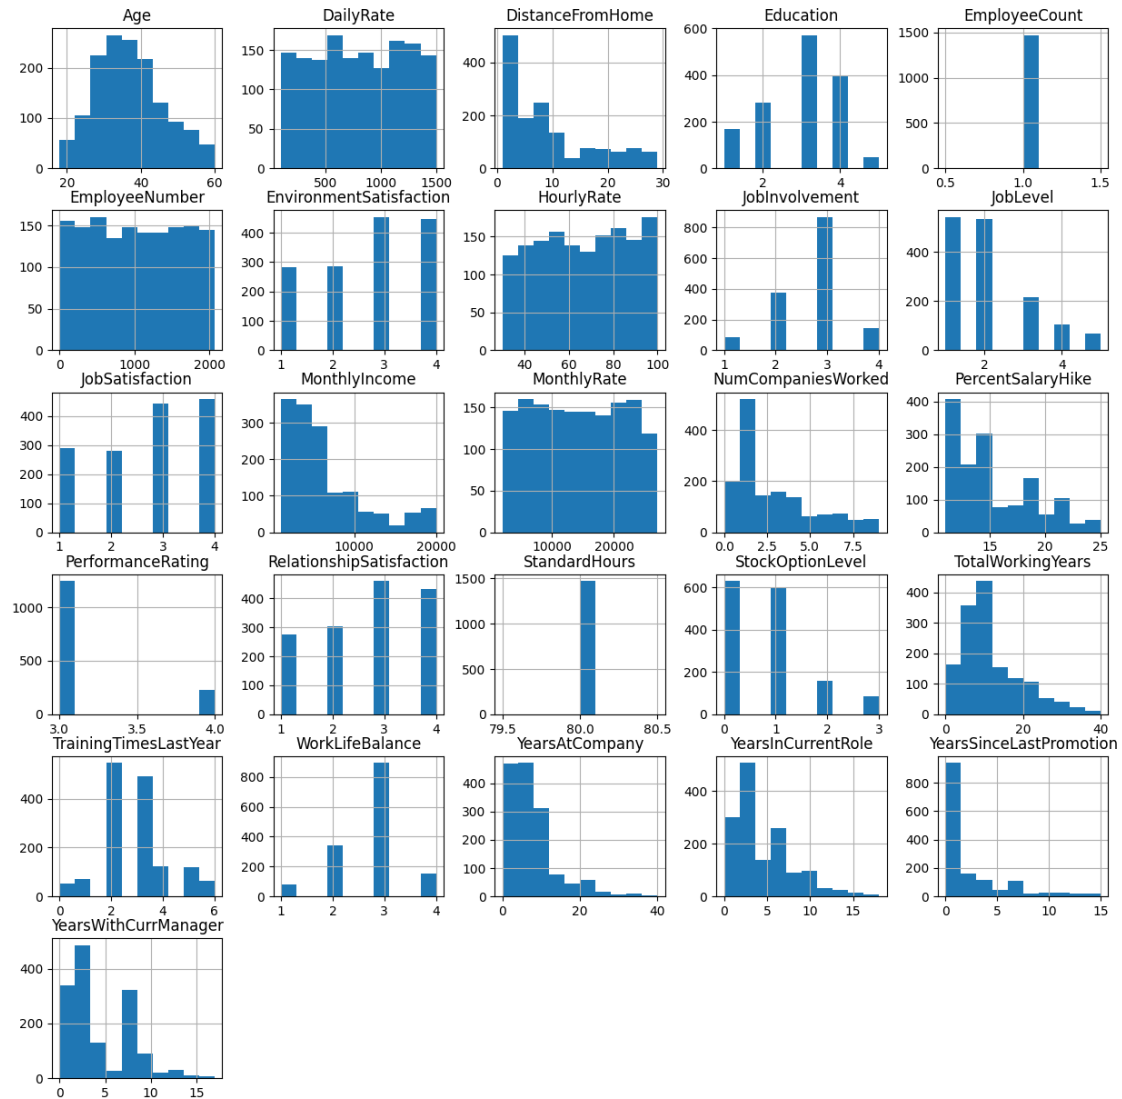
```
[7]: df.hist(figsize=(15,15))
```

```
[7]: array([[<Axes: title={'center': 'Age'}>,
          <Axes: title={'center': 'DailyRate'}>],
```

```

<Axes: title={'center': 'DistanceFromHome'}>,
<Axes: title={'center': 'Education'}>,
<Axes: title={'center': 'EmployeeCount'}>],
[<Axes: title={'center': 'EmployeeNumber'}>,
<Axes: title={'center': 'EnvironmentSatisfaction'}>,
<Axes: title={'center': 'HourlyRate'}>,
<Axes: title={'center': 'JobInvolvement'}>,
<Axes: title={'center': 'JobLevel'}>],
[<Axes: title={'center': 'JobSatisfaction'}>,
<Axes: title={'center': 'MonthlyIncome'}>,
<Axes: title={'center': 'MonthlyRate'}>,
<Axes: title={'center': 'NumCompaniesWorked'}>,
<Axes: title={'center': 'PercentSalaryHike'}>],
[<Axes: title={'center': 'PerformanceRating'}>,
<Axes: title={'center': 'RelationshipSatisfaction'}>,
<Axes: title={'center': 'StandardHours'}>,
<Axes: title={'center': 'StockOptionLevel'}>,
<Axes: title={'center': 'TotalWorkingYears'}>],
[<Axes: title={'center': 'TrainingTimesLastYear'}>,
<Axes: title={'center': 'WorkLifeBalance'}>,
<Axes: title={'center': 'YearsAtCompany'}>,
<Axes: title={'center': 'YearsInCurrentRole'}>,
<Axes: title={'center': 'YearsSinceLastPromotion'}>],
[<Axes: title={'center': 'YearsWithCurrManager'}>, <Axes: >,
<Axes: >, <Axes: >, <Axes: >]], dtype=object)


```



```
[8]: #inference: Monthly Income, Total Working hours, Years at company, Distance
      ↪from home are all right skewed.
      #inference: Employee count and Standard hours are redundant so they can be
      ↪removed.
```

```
[9]: sns.
      ↪catplot(x='JobRole',hue='Attrition',data=df,kind='count',height=7,aspect=3,legend=False)
```

```
[9]: <seaborn.axisgrid.FacetGrid at 0x7ddce526b280>
```




```
[10]: corr=df.corr()
```

<ipython-input-10-0014364bc22a>:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
corr=df.corr()
```

```
[11]: plt.figure(figsize=(8, 6))
sns.countplot(x='Attrition', data=df)
plt.title('Attrition Count')
plt.xlabel('Attrition')
plt.ylabel('Count')
plt.show()
```

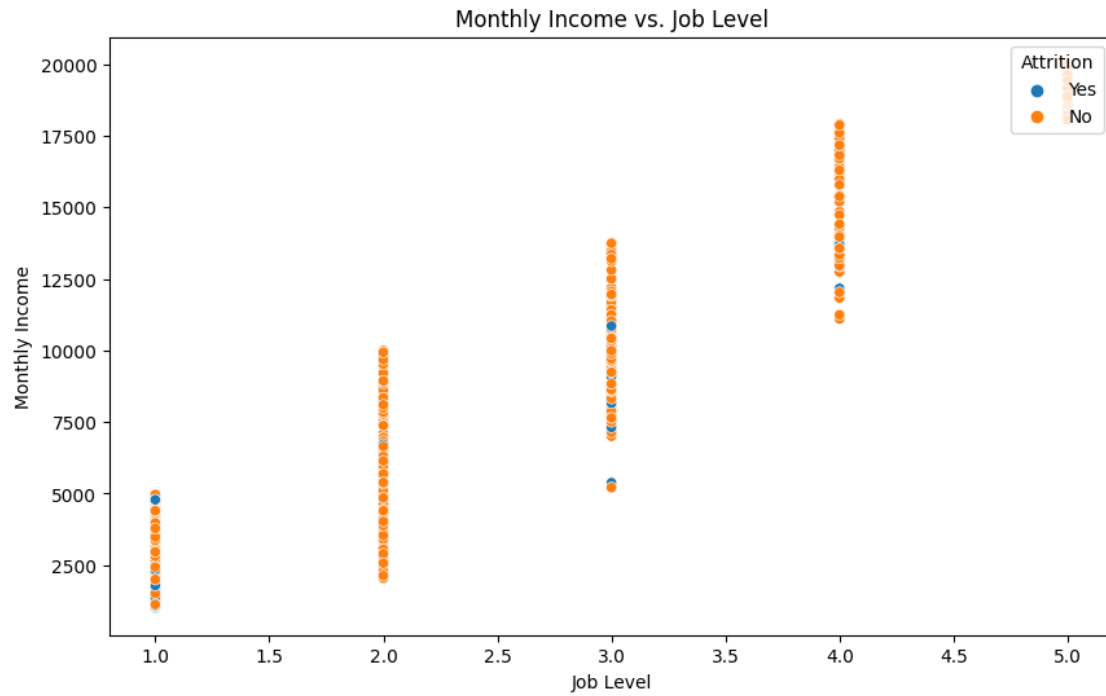




```
[12]: #inference: The plot suggests that the attrition is relatively less common in  
      ↪ dataset
```

```
[13]: plt.figure(figsize=(10, 6))  
      sns.scatterplot(x='JobLevel', y='MonthlyIncome', hue='Attrition', data=df)  
      plt.title('Monthly Income vs. Job Level')  
      plt.xlabel('Job Level')  
      plt.ylabel('Monthly Income')  
      plt.legend(title='Attrition', loc='upper right')
```


```
[13]: <matplotlib.legend.Legend at 0x7ddca1c91c90>
```




```
[14]: #inference: As the job level increases the monthly income also increases
```

```
[15]: sns.heatmap(corr,vmax=0.5,linewidth=0.2)
```

```
[15]: <Axes: >
```



```
[16]: plt.figure(figsize=(8, 6))
sns.boxplot(x=df['MonthlyIncome'])
plt.xlabel('')
plt.show()
```



```
[17]: Q1 = df['MonthlyIncome'].quantile(0.25)
      Q3 = df['MonthlyIncome'].quantile(0.75)
      IQR = Q3 - Q1
      outliers = (df['MonthlyIncome'] < Q1 - 1.5 * IQR) | (df['MonthlyIncome'] > Q3 +
      ↪ 1.5 * IQR)
      outlier_rows = df[outliers]
      print(outlier_rows)
```

	Age	Attrition	BusinessTravel	DailyRate	Department \
25	53	No	Travel_Rarely	1282	Research & Development
29	46	No	Travel_Rarely	705	Sales
45	41	Yes	Travel_Rarely	1360	Research & Development
62	50	No	Travel_Rarely	989	Research & Development
105	59	No	Non-Travel	1420	Human Resources
...	...	...	...	...	...
1374	58	No	Travel_Rarely	605	Sales
1377	49	No	Travel_Frequently	1064	Research & Development
1401	55	No	Travel_Rarely	189	Human Resources
1437	39	No	Non-Travel	105	Research & Development
1443	42	No	Travel_Rarely	300	Research & Development

	DistanceFromHome	Education	EducationField	EmployeeCount	\
25	5	3	Other	1	
29	2	4	Marketing	1	
45	12	3	Technical Degree	1	
62	7	2	Medical	1	
105	2	4	Human Resources	1	
...	...	...	...	...	
1374	21	3	Life Sciences	1	
1377	2	1	Life Sciences	1	
1401	26	4	Human Resources	1	
1437	9	3	Life Sciences	1	
1443	2	3	Life Sciences	1	

	EmployeeNumber	...	RelationshipSatisfaction	StandardHours	\
25	32	...	4	80	
29	38	...	4	80	
45	58	...	4	80	
62	80	...	4	80	
105	140	...	4	80	
...	...	...	...	...	
1374	1938	...	3	80	
1377	1941	...	4	80	
1401	1973	...	1	80	
1437	2022	...	3	80	
1443	2031	...	1	80	

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
25	1	26	3	
29	0	22	2	
45	0	23	0	
62	1	29	2	
105	1	30	3	
...	...	...	...	
1374	1	29	2	
1377	0	28	3	
1401	1	35	0	
1437	0	21	3	
1443	0	24	2	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
25	2	14	13	
29	2	2	2	
45	3	22	15	
62	2	27	3	
105	3	3	2	
...	...	...	...	
1374	2	1	0	

1377	3	5	4
1401	3	10	9
1437	2	6	0
1443	2	22	6

	YearsSinceLastPromotion	YearsWithCurrManager
25	4	8
29	2	1
45	15	8
62	13	8
105	2	2
...	...	...
1374	0	0
1377	4	3
1401	1	4
1437	1	3
1443	4	14

[114 rows x 35 columns]

```
[18]: df= pd.get_dummies(df, columns=['BusinessTravel', 'Department',
↳ 'EducationField', 'Gender', 'JobRole', 'MaritalStatus', 'Over18',
↳ 'OverTime'])
```

```
[19]: from sklearn.preprocessing import LabelEncoder
lb=LabelEncoder()
df['Attrition']=lb.fit_transform(df['Attrition'])
```

```
[20]: x=df.drop('Attrition',axis=1)
y=df['Attrition']
```

```
[21]: x.head()
```

```
[21]:   Age  DailyRate  DistanceFromHome  Education  EmployeeCount  EmployeeNumber  \
0   41      1102             1         2           1             1
1   49       279             8         1           1             2
2   37      1373             2         2           1             4
3   33      1392             3         4           1             5
4   27       591             2         1           1             7
```

	EnvironmentSatisfaction	HourlyRate	JobInvolvement	JobLevel	...	\
0	2	94	3	2	...	
1	3	61	2	2	...	
2	4	92	2	1	...	
3	4	56	3	1	...	
4	1	40	3	1	...	

	JobRole_Research Director	JobRole_Research Scientist \
0	0	0
1	0	1
2	0	0
3	0	1
4	0	0

	JobRole_Sales Executive	JobRole_Sales Representative \
0	1	0
1	0	0
2	0	0
3	0	0
4	0	0

	MaritalStatus_Divorced	MaritalStatus_Married	MaritalStatus_Single \
0	0	0	1
1	0	1	0
2	0	0	1
3	0	1	0
4	0	1	0

	Over18_Y	OverTime_No	OverTime_Yes
0	1	0	1
1	1	1	0
2	1	0	1
3	1	0	1
4	1	1	0

[5 rows x 55 columns]

```
[22]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
```

```
[23]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
[23]: ((1176, 55), (294, 55), (1176,), (294,))
```

```
[24]: from sklearn.preprocessing import MinMaxScaler
sc=MinMaxScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.transform(x_test)
x_train,x_test
```

```
[24]: (array([[0.95238095, 0.35913978, 0.71428571, ..., 0.          , 0.          ,
1.          ],
[0.64285714, 0.60645161, 0.96428571, ..., 0.          , 1.          ,
0.          ],
```

```

[0.52380952, 0.14050179, 0.89285714, ..., 0.          , 1.          ,
 0.          ],
...,
[0.5952381 , 0.77060932, 0.03571429, ..., 0.          , 0.          ,
 1.          ],
[0.47619048, 0.11756272, 0.03571429, ..., 0.          , 0.          ,
 1.          ],
[0.52380952, 0.39713262, 0.32142857, ..., 0.          , 1.          ,
 0.          ]]),
array([[0.42857143, 0.38064516, 0.32142857, ..., 0.          , 1.          ,
 0.          ],
[0.35714286, 0.33763441, 0.85714286, ..., 0.          , 1.          ,
 0.          ],
[0.4047619 , 0.4          , 0.60714286, ..., 0.          , 0.          ,
 1.          ],
...,
[0.30952381, 0.10394265, 0.17857143, ..., 0.          , 0.          ,
 1.          ],
[0.47619048, 0.82939068, 0.03571429, ..., 0.          , 1.          ,
 0.          ],
[0.52380952, 0.18996416, 0.25          , ..., 0.          , 0.          ,
 1.          ]]))

```

```

[45]: from sklearn.linear_model import LogisticRegression
x_train_logistic=x_train
y_train_logistic=y_train
x_test_logistic=x_test
y_test_logistic=y_test
model=LogisticRegression()
model.fit(x_train_logistic,y_train_logistic)

```

[45]: LogisticRegression()

```

[46]: pred_logistic=model.predict(x_test_logistic)
pred_logistic

```

```

[46]: array([0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,

```



```
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0])
```

```
[47]: y_test_logistic
```

```
[47]: 442      0
      1091    0
      981     1
      785     0
      1332    1
      ..
      1439    0
      481     0
      124     1
      198     0
      1229    0
      Name: Attrition, Length: 294, dtype: int64
```

```
[48]: from sklearn.metrics import
      accuracy_score, confusion_matrix, classification_report, roc_auc_score, roc_curve, precision_score
      accuracy_logistic = accuracy_score(y_test_logistic, pred_logistic)
      print("Accuracy:", accuracy_logistic)
      precision_logistic = precision_score(y_test_logistic, pred_logistic)
      print("Precision:", precision_logistic)
      recall_logistic = recall_score(y_test_logistic, pred_logistic)
      print("Recall:", recall_logistic)
      f1_logistic = f1_score(y_test_logistic, pred_logistic)
      print("F1 Score:", f1_logistic)
      roc_auc_logistic = roc_auc_score(y_test_logistic, pred_logistic)
      print("ROC AUC Score:", roc_auc_logistic)
```

```
Accuracy: 0.8877551020408163
Precision: 0.8076923076923077
Recall: 0.42857142857142855
F1 Score: 0.5599999999999999
ROC AUC Score: 0.7040816326530612
```

```
[49]: print(classification_report(y_test_logistic, pred_logistic))
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	245
1	0.81	0.43	0.56	49
accuracy			0.89	294
macro avg	0.85	0.70	0.75	294

weighted avg      0.88      0.89      0.87      294

```
[50]: from sklearn.tree import DecisionTreeClassifier
      x_train_dtc=x_train
      y_train_dtc=y_train
      x_test_dtc=x_test
      y_test_dtc=y_test
      dtc=DecisionTreeClassifier()
      dtc.fit(x_train_dtc,y_train_dtc)
```

```
[50]: DecisionTreeClassifier()
```

```
[52]: pred_dtc=dtc.predict(x_test_dtc)
      pred_dtc
```

```
[52]: array([0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
          0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0,
          0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1,
          1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0,
          0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0,
          0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1,
          0, 1, 0, 0, 0, 0, 0, 0])
```

```
[53]: y_test_dtc
```

```
[53]: 442      0
      1091     0
      981      1
      785      0
      1332     1
      ..
      1439     0
      481      0
      124      1
      198      0
      1229     0
      Name: Attrition, Length: 294, dtype: int64
```

```
[54]: accuracy_dtc = accuracy_score(y_test_dtc, pred_dtc)
print("Accuracy:", accuracy_dtc)
precision_dtc = precision_score(y_test_dtc, pred_dtc)
print("Precision:", precision_dtc)
recall_dtc = recall_score(y_test_dtc, pred_dtc)
print("Recall:", recall_dtc)
f1_dtc = f1_score(y_test_dtc, pred_dtc)
print("F1 Score:", f1_dtc)
roc_auc_dtc = roc_auc_score(y_test_dtc, pred_dtc)
print("ROC AUC Score:", roc_auc_dtc)
```

```
Accuracy: 0.7755102040816326
Precision: 0.3333333333333333
Recall: 0.3469387755102041
F1 Score: 0.33999999999999997
ROC AUC Score: 0.6040816326530611
```

```
[55]: print(classification_report(y_test_dtc,pred_dtc))
```

	precision	recall	f1-score	support
0	0.87	0.86	0.86	245
1	0.33	0.35	0.34	49
accuracy			0.78	294
macro avg	0.60	0.60	0.60	294
weighted avg	0.78	0.78	0.78	294

```
[35]: from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV
```

```
[36]: param_grid = {
    'solver':['newton-cg','liblinear'],
    'penalty':['l2'],
    'C':np.logspace(-4.5,4.5,50),
    'class_weight':['balanced'],
    'tol':[0.0001,0.001,0.01,0.1],
    'fit_intercept':[True,False],
    'intercept_scaling':[1,2,3]
}
```

```
[56]: grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=5,
    ↪scoring='roc_auc', n_jobs=-1)
grid_search.fit(x_train_logistic, y_train_logistic)
```

```
[56]: GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,
                param_grid={'C': array([3.16227766e-05, 4.82695744e-05,
7.36795456e-05, 1.12465782e-04,
1.71669791e-04, 2.62039853e-04, 3.99982340e-04, 6.10540230e-04,
9.31939576e-04, 1.42252931e-03, 2.17137430e-03, 3.31442475e-03,
5.05919749e-03, 7.72244995e-03, 1.17876863e-02, 1.79929362e-02,
2.74647411e-02, 4.19226744e-02,
1.29492584e+02, 1.97659807e+02, 3.01711481e+02, 4.60537826e+02,
7.02973212e+02, 1.07303094e+03, 1.63789371e+03, 2.50011038e+03,
3.81621341e+03, 5.82513671e+03, 8.89159334e+03, 1.35722878e+04,
2.07169840e+04, 3.16227766e+04]),
                'class_weight': ['balanced'],
                'fit_intercept': [True, False],
                'intercept_scaling': [1, 2, 3], 'penalty': ['l2'],
                'solver': ['newton-cg', 'liblinear'],
                'tol': [0.0001, 0.001, 0.01, 0.1]},
                scoring='roc_auc')
```

```
[57]: grid_search.fit(x_train_logistic,y_train_logistic)
```

```
[57]: GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,
                param_grid={'C': array([3.16227766e-05, 4.82695744e-05,
7.36795456e-05, 1.12465782e-04,
1.71669791e-04, 2.62039853e-04, 3.99982340e-04, 6.10540230e-04,
9.31939576e-04, 1.42252931e-03, 2.17137430e-03, 3.31442475e-03,
5.05919749e-03, 7.72244995e-03, 1.17876863e-02, 1.79929362e-02,
2.74647411e-02, 4.19226744e-02,
1.29492584e+02, 1.97659807e+02, 3.01711481e+02, 4.60537826e+02,
7.02973212e+02, 1.07303094e+03, 1.63789371e+03, 2.50011038e+03,
3.81621341e+03, 5.82513671e+03, 8.89159334e+03, 1.35722878e+04,
2.07169840e+04, 3.16227766e+04]),
                'class_weight': ['balanced'],
                'fit_intercept': [True, False],
                'intercept_scaling': [1, 2, 3], 'penalty': ['l2'],
                'solver': ['newton-cg', 'liblinear'],
                'tol': [0.0001, 0.001, 0.01, 0.1]},
                scoring='roc_auc')
```

```
[58]: best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
```

```
[59]: print(classification_report(y_test_logistic,pred_logistic))
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	245
1	0.81	0.43	0.56	49

accuracy			0.89	294
macro avg	0.85	0.70	0.75	294
weighted avg	0.88	0.89	0.87	294

```
[61]: accuracy_logistic = accuracy_score(y_test_logistic, pred_logistic)
print("Accuracy:", accuracy_logistic)
precision_logistic = precision_score(y_test_logistic, pred_logistic)
print("Precision:", precision_logistic)
recall_logistic = recall_score(y_test_logistic, pred_logistic)
print("Recall:", recall_logistic)
f1_logistic = f1_score(y_test_logistic, pred_logistic)
print("F1 Score:", f1_logistic)
roc_auc_logistic = roc_auc_score(y_test_logistic, pred_logistic)
print("ROC AUC Score:", roc_auc_logistic)
```

```
Accuracy: 0.8877551020408163
Precision: 0.8076923076923077
Recall: 0.42857142857142855
F1 Score: 0.5599999999999999
ROC AUC Score: 0.7040816326530612
```

```
[42]: #inference: There is a class imbalance-Attrition=1
```